

Using AWS From the
Command Line Interface

AWS CLI CHEAT SHEET



AWS CLI Cheat Sheet

Table of Contents

- [Volumes](#)
 - [Describing volumes](#)
 - [Describing volumes using a different aws user profile](#)
 - [Listing Available Volumes IDs](#)
 - [Deleting a Volume](#)
 - [Deleting Unused Volumes.. Think Before You Type :-\)](#)
 - [Creating a Snapshot](#)
 - [Creating an Image \(AMI\)](#)
 - [Creating AMI Without Rebooting the Machine](#)
- [AMIs](#)
 - [Listing AMI\(s\)](#)
 - [Describing AMI\(s\)](#)
 - [Listing Amazon AMIs](#)
 - [Using Filters](#)
- [Lambda](#)
 - [List Functions](#)
 - [Describe a Function](#)
 - [Invoke a Function](#)
 - [Update a Function Code](#)
 - [Publish a Version](#)
 - [List Layers](#)
 - [List Aliases of a Function](#)
 - [Describe an Alias](#)
 - [Create an Alias](#)
 - [Delete an Alias](#)
 - [List Function Tags](#)
 - [Delete a Function](#)
 - [Using AWS Lambda with Scheduled Events](#)
- [IAM](#)
 - [List Users](#)
 - [List Policies](#)
 - [List Groups](#)
 - [Get Users in a Group](#)
 - [Describing a Policy.](#)
 - [List Access Keys](#)
 - [List Keys](#)
 - [List the Access Key IDs for an IAM User](#)
 - [List the SSH Public Keys for a User](#)
- [S3 API](#)
 - [Listing Buckets](#)
 - [Listing Only Bucket Names](#)

- [Getting a Bucket Region](#)
- [Listing the Content of a Bucket](#)
- [Syncing a Local Folder with a Bucket](#)
- [Copying Files](#)
- [Copying Folders](#)
- [Removing a File from a Bucket](#)
- [Deleting a Bucket](#)
- [Emptying a Bucket](#)
- [VPC](#)
 - [Creating A VPC](#)
 - [Allowing DNS hostnames](#)
- [Subnets](#)
 - [Creating A Subnet](#)
 - [Auto Assigning Public IPs To Instances In A Public Subnet](#)
- [Internet Gateway](#)
 - [Creating An IGW](#)
 - [Attaching An IGW to A VPC](#)
- [NAT](#)
 - [Setting Up A NAT Gateway](#)
- [Route Tables](#)
 - [Creating A Public Route Table](#)
 - [Creating A Private Route Tables](#)
- [CloudFront](#)
 - [Listing Distributions](#)
 - [Invalidating Files From a Distribution](#)
 - [Sync a Local Folder with a CloudFront Distribution](#)
 - [Sync and Invalidate at the Same Time:](#)
- [RDS](#)
 - [List Databases](#)
 - [List Public Databases](#)
 - [List Non Protected Databases \(DeletionProtection\)](#)
 - [Describe the Automated Backups for a DB Instance](#)
 - [Create a DB Cluster](#)
 - [Create a DB Instance](#)
 - [Create a DB Security Group](#)
 - [Create a Read Replica](#)
 - [Create a Custom DB Cluster Endpoint](#)
 - [Apply Tag to a DB](#)
 - [Create a Cluster Snapshot](#)
 - [Create a CloudWatch Alarm for a DB Instance](#)
- [Connect Deeper](#)

Volumes

Describing volumes

```
aws ec2 describe-volumes
```

Describing filtered volumes:

```
aws ec2 describe-volumes --filters Name=status,Values=creating | available |  
in-use | deleting | deleted | error
```

e.g, describing all deleted volumes:

```
aws ec2 describe-volumes --filters Name=status,Values=deleted
```

Filters can be applied to the attachment status:

```
aws ec2 describe-volumes --filters Name=attachment.status,Values=attaching |  
attached | detaching | detached
```

e.g: describing all volumes with the status "attaching":

```
aws ec2 describe-volumes --filters Name=attachment.status,Values=attaching
```

This is the generic form. Use `--profile <your_profile_name>`, if you have multiple AWS profiles or accounts.

```
aws ec2 describe-volumes --filters Name:'tag:Name',Values: ['some_values'] --  
profile <your_profile_name>
```

Describing volumes using a different aws user profile

```
aws ec2 describe-volumes --filters Name=status,Values=in-use --profile  
<your_profile_name>
```

Listing Available Volumes IDs

```
aws ec2 describe-volumes --filters Name=status,Values=available |grep  
VolumeId|awk '{print $2}' | tr '\n|,' ' ' ' '
```

With "profile":

```
aws ec2 describe-volumes --filters Name=status,Values=available --profile  
<your_profile_name>|grep VolumeId|awk '{print $2}' | tr '\n|,' ' ' ' '
```

Deleting a Volume

```
aws ec2 delete-volume --region <region> --volume-id <volume_id>
```

Deleting Unused Volumes.. Think Before You Type :-)

```
for x in $(aws ec2 describe-volumes --filters Name=status,Values=available --profile <your_profile_name>|grep VolumeId|awk '{print $2}' | tr ',|"' ' '); do  
aws ec2 delete-volume --region <region> --volume-id $x; done
```

With "profile":

```
for x in $(aws ec2 describe-volumes --filters Name=status,Values=available --profile <your_profile_name>|grep VolumeId|awk '{print $2}' | tr ',|"' ' '); do  
aws ec2 delete-volume --region <region> --volume-id $x --profile  
<your_profile_name>; done
```

Creating a Snapshot

```
aws ec2 create-snapshot --volume-id <vol-id>
```

```
aws ec2 create-snapshot --volume-id <vol-id> --description "snapshot-$(date  
+'%Y-%m-%d_%H-%M-%S')"
```

Creating an Image (AMI)

```
aws ec2 create-image --instance-id <instance_id> --name "image-$(date +%Y-%m-%d_%H-%M-%S')" --description "image-$(date +%Y-%m-%d_%H-%M-%S')"
```

Creating AMI Without Rebooting the Machine

```
aws ec2 create-image --instance-id <instance_id> --name "image-$(date +%Y-%m-%d_%H-%M-%S')" --description "image-$(date +%Y-%m-%d_%H-%M-%S')"
```

You are free to change the AMI name `image-$(date +%Y-%m-%d_%H-%M-%S')` to a name of your choice.

AMIs

Listing AMI(s)

```
aws ec2 describe-images
```

Describing AMI(s)

```
aws ec2 describe-images --image-ids <image_id> --profile <profile> --region  
<region>
```

e.g:

```
aws ec2 describe-images --image-ids ami-e24dfa9f --profile terraform --region  
eu-west-3
```

Listing Amazon AMIs

```
aws ec2 describe-images --owners amazon
```

Using Filters

e.g: Describing Windows AMIs that are backed by Amazon EBS.

```
aws ec2 describe-images --filters "Name=platform,Values=windows" "Name=root-  
device-type,Values=ebs"
```

e.g: Describing Ubuntu AMIs

```
aws ec2 describe-images --filters "Name=name,Values=ubuntu*"
```

Lambda

List Functions

```
aws lambda list-functions
```

Describe a Function

```
aws lambda get-function --function-name my-function
```

Invoke a Function

```
aws lambda invoke --function-name my-function --payload '{ "name": "Bob" }'
response.json
```

Update a Function Code

```
aws lambda update-function-code --function-name my-function --zip-file
fileb://my-function.zip
```

Publish a Version

```
aws lambda publish-version --function-name my-function
```

List Layers

Let's take this example in which we want to list information of layers that are compatible with Python 3.7 runtime.

```
aws lambda list-layers --compatible-runtime python3.7
```

Possible layers runtime:

```
nodejs
nodejs4.3
nodejs6.10
nodejs8.10
nodejs10.x
nodejs12.x
java8
java8.a12
java11
python2.7
python3.6
python3.7
python3.8
dotnetcore1.0
dotnetcore2.0
dotnetcore2.1
dotnetcore3.1
nodejs4.3-edge
```

```
go1.x
ruby2.5
ruby2.7
provided
provided.al2
```

List Aliases of a Function

```
aws lambda list-aliases --function-name my-function
```

Describe an Alias

```
aws lambda get-alias --function-name my-function --name LIVE
```

Create an Alias

```
aws lambda create-alias --function-name my-function --description "alias
description goes here" --function-version 1 --name LIVE
```

Delete an Alias

```
aws lambda delete-alias --function-name my-function --name LIVE
```

List Function Tags

```
aws lambda list-tags --resource arn:aws:lambda:eu-west-
1:xxxxxxxxxxxx:function:my-function
```

Delete a Function

```
aws lambda delete-function --function-name my-function
```

Using AWS Lambda with Scheduled Events

```
sid=Sid$(date +%Y%m%d%H%M%S); aws lambda add-permission --statement-id $sid --  
action 'lambda:InvokeFunction' --principal events.amazonaws.com --source-arn  
arn:aws:events:<region>:<arn>:rule/AWSLambdaBasicExecutionRole --function-name  
function:<awsents> --region <region>
```

##

IAM

List Users

```
aws iam list-users
```

List Policies

```
aws iam list-policies
```

List Groups

```
aws iam list-groups
```

Get Users in a Group

```
aws iam get-group --group-name <group_name>
```

Describing a Policy

```
aws iam get-policy --policy-arn arn:aws:iam::aws:policy/<policy_name>
```

List Access Keys

```
aws iam list-access-keys
```

List Keys

```
aws iam list-access-keys
```

List the Access Key IDs for an IAM User

```
aws iam list-access-keys --user-name <user_name>
```

List the SSH Public Keys for a User

```
aws iam list-ssh-public-keys --user-name <user_name>
```

S3 API

Listing Buckets

```
aws s3api list-buckets
```

Or

```
aws s3 ls
```

e.g

```
aws s3 ls --profile eon01
```

Listing Only Bucket Names

```
aws s3api list-buckets --query 'Buckets[].Name'
```

Getting a Bucket Region

```
aws s3api get-bucket-location --bucket <bucket_name>
```

e.g

```
aws s3api get-bucket-location --bucket practicalaws.com
```

Listing the Content of a Bucket

```
aws s3 ls s3://<bucket_name> --region <region>
```

e.g

```
aws s3 ls s3://practicalaws.com
```

```
aws s3 ls s3://practicalaws.com --region eu-west-1
```

```
aws s3 ls s3://practicalaws.com --region eu-west-1 --profile eon01
```

Syncing a Local Folder with a Bucket

```
aws s3 sync <local_path> s3://<bucket_name>
```

e.g

```
aws s3 sync . s3://practicalaws.com --region eu-west-1
```

Copying Files

```
aws s3 cp <file_name> s3://<bucket_name>
```

Or:

```
aws s3 cp <file_name> s3://<bucket_name>/<folder_name>/
```

To copy all files from a folder, look at "Copying Folders". Or use the following example, where I copy the content of the folder "images (contains images)" in the remote folder "images".

```
cd images
```

```
aws s3 cp . s3://saltstackfordevops.com/images --recursive --region us-east-2
```

Copying Folders

```
aws s3 cp <folder_name>/ s3://<bucket_name>/ --recursive
```

To exclude files:

```
aws s3 cp <folder_name>/ s3://<bucket_name>/ --recursive --exclude "  
<file_name_or_a_wildcard>"
```

e.g: To only include a certain type of files (PNG) and exclude others (JPG)

```
aws s3 cp practicalaws.com/ s3://practicalaws-backup/ --recursive --exclude  
"*.jpg" --include "*.png"
```

e.g: To exclude a folder

```
aws s3 cp practicalaws.com/ s3://practicalaws-backup/ --recursive --exclude  
".git/*"
```

Removing a File from a Bucket

```
aws s3 rm s3://<bucket_name>/<object_name>
```

e.g

```
aws s3 rm s3://practicalaws.com/temp.txt
```

Deleting a Bucket

```
aws s3 rb s3://<bucket_name> --force
```

If the bucket is not empty, use --force.

e.g

```
aws s3 rb s3://practicalaws.com --force
```

Emptying a Bucket

```
aws s3 rm s3://<bucket_name>/<key_name> --recursive
```

e.g

In order to remove tempfiles/file1.txt and tempfiles/file2.txt from practicalaws.com bucket, use:

```
aws s3 rm s3://practicalaws.com/tempfiles --recursive
```

Remove all objects using:

```
aws s3 rm s3://practicalaws.com/tempfiles
```

Making a Public File Private

```
aws s3api put-object-acl --acl private --bucket <bucket-name> --key <file_name  
or file_path>
```

e.g:

```
aws s3api put-object-acl --acl private --bucket practicalaws.com --key  
image/logo.png
```

Making a Public bucket Private

```
aws s3 ls --recursive s3://<bucket-name> | cut -d' ' -f5- | awk '{print $NF}' |  
while read line; do  
    echo "$line"  
    aws s3api put-object-acl --acl private --bucket <bucket-name> --key "$line"  
done
```

VPC

Creating A VPC

```
aws ec2 create-vpc --cidr-block <cidr_block> --regionsn <region>
```

e.g

```
aws ec2 create-vpc --cidr-block 10.0.0.0/16 --region eu-west-1
```

Allowing DNS hostnames

```
aws ec2 modify-vpc-attribute --vpc-id <vpc_id> --enable-dns-hostnames "  
{\"Value\":true}\" --region <region>
```

Subnets

Creating A Subnet

```
aws ec2 create-subnet --vpc-id <vpc_id> --cidr-block <cidr_block> --availability-zone <availability_zone> --region <region>
```

Auto Assigning Public IPs To Instances In A Public Subnet

```
aws ec2 modify-subnet-attribute --subnet-id <subnet_id> --map-public-ip-on-launch --region <region>
```

Internet Gateway

Creating An IGW

```
aws ec2 create-internet-gateway --region <region>
```

Attaching An IGW to A VPC

```
aws ec2 attach-internet-gateway --internet-gateway-id <igw_id> --vpc-id <vpc_id> --region <region>
```

NAT

Setting Up A NAT Gateway

Allocate Elastic IP

```
aws ec2 allocate-address --domain vpc --region <region>
```

then use the AllocationId to create the NAT Gateway for the public zone in

```
aws ec2 create-nat-gateway --subnet-id <subnet_id> --allocation-id <allocation_id> --region <region>
```

Route Tables

Creating A Public Route Table

Create the Route Table:

```
aws ec2 create-route-table --vpc-id <vpc_id> --region <region>
```

then create a route for an Internet Gateway.

Now, use the outputted Route Table ID:

```
aws ec2 create-route --route-table-id <route_table_id> --destination-cidr-block 0.0.0.0/0 --gateway-id <igw_id> --region <region>
```

Finally, associate the public subnet with the Route Table

```
aws ec2 associate-route-table --route-table-id <route_table_id> --subnet-id <subnet_id> --region <region>
```

Creating A Private Route Tables

Create the Route Table

```
aws ec2 create-route-table --vpc-id <vpc_id> --region <region>
```

then create a route that points to a NAT Gateway

```
aws ec2 create-route --route-table-id <route_table_id> --destination-cidr-block 0.0.0.0/0 --nat-gateway-id <net_gateway_id> --region <region>
```

Finally, associate the subnet

```
aws ec2 associate-route-table --route-table-id <route_table_id> --subnet-id <subnet_id> --region <region>
```

CloudFront

Listing Distributions

In some cases, you need to setup this first:

```
aws configure set preview.cloudfront true
```

Then:

```
aws cloudfront list-distributions
```

Invalidating Files From a Distribution

To invalidate index and error HTML files from the distribution with the ID Z2W2LX9VBMAPRX:

```
aws cloudfront create-invalidation --distribution-id Z2W2LX9VBMAPRX --paths /index.html /error.html
```

To invalidate everything in the distribution:

```
aws cloudfront create-invalidation --distribution-id Z2W2LX9VBMAPRX --paths '/*'
```

Sync a Local Folder with a CCloudFront Distribution

CloudFront is "attached" to a bucket, you need to upload your files to the bucket.

e.g.:

```
aws s3 sync . s3://my-bucket.com
```

If you should keep the files public:

```
aws s3 sync . s3://my-bucket.com --acl public-read
```

To copy a single file, you need to:

```
aws s3 cp file1 s3://my-bucket.com/sub-folder/ --acl <ACL>
```

Sync and Invalidate at the Same Time:

```
aws s3 sync . s3://my-bucket.com --acl public-read && aws cloudfront create-invalidation --distribution-id Z2W2LX9VBMAPRX --paths '/*'
```

or in case you want to update a single file:

```
aws s3 cp file1 s3://my-bucket.com/sub-folder/ --acl public-read && aws cloudfront create-invalidation --distribution-id Z2W2LX9VBMAPRX --paths '/sub-folder/file1'
```

RDS

List Databases

```
aws rds describe-db-instances
```

or:

```
aws rds describe-db-instances --query 'DBInstances[].DBInstanceIdentifier'
```

List Public Databases

```
aws rds describe-db-instances --query 'DBInstances[?PubliclyAccessible=="true"].
[DBInstanceIdentifier,Endpoint.Address]'
```

List Non Protected Databases (DeletionProtection)

```
aws rds describe-db-instances \
  --query 'DBInstances[*].[DBInstanceIdentifier]' \
  --output text \
  | xargs -I {} bash -c 'if [[ $(aws rds describe-db-instances --db-instance-
  identifier {} --query "'DBInstances[*].DeletionProtection'" --output text)
  == False ]]; then echo {} ; fi'
```

Describe the Automated Backups for a DB Instance

```
aws rds describe-db-instance-automated-backups --db-instance-identifier
database-mysql
```

Create a DB Cluster

```
aws rds create-db-cluster \
  --db-cluster-identifier mysql-cluster \
  --engine aurora-mysql \
  --engine-version 5.7.12 \
  --master-username master \
  --master-user-password xxxxxx \
  --db-subnet-group-name default \
  --vpc-security-group-ids sg-0130572b9daf3dc16
```

Create a DB Instance

```
aws rds create-db-instance \  
  --db-instance-identifier mysql-instance \  
  --db-instance-class db.t3.micro \  
  --engine mysql \  
  --master-username admin \  
  --master-user-password xxxxx \  
  --allocated-storage 40
```

Create a DB Security Group

```
aws rds create-db-security-group --db-security-group-name my-security-group --  
db-security-group-description "My Security Group"
```

Create a Read Replica

```
aws rds create-db-instance-read-replica \  
  --db-instance-identifier test-instance-repl \  
  --source-db-instance-identifier test-instance
```

Create a Custom DB Cluster Endpoint

```
aws rds create-db-cluster-endpoint \  
  --db-cluster-endpoint-identifier mycustomendpoint \  
  --endpoint-type reader \  
  --db-cluster-identifier mydbcluster \  
  --static-members dbinstance1 dbinstance2
```

Apply Tag to a DB

```
aws rds add-tags-to-resource \  
  --resource-name arn:aws:rds:us-east-1:123456789012:db:database-mysql \  
  --tags "[{\"Key\": \"Name\", \"Value\": \"MyDatabase\"}, {\"Key\":  
  \"Environment\", \"Value\": \"test\"}]"
```

Create a Cluster Snapshot

```
aws rds create-db-cluster-snapshot --db-cluster-identifier my-db-cluster --db-  
cluster-snapshot-identifier my-db-cluster-snapshot
```

Create a CloudWatch Alarm for a DB Instance

e.g.: When **average CPU for latest 15 minutes is above 90%**

```
aws cloudwatch put-metric-alarm \  
  --alarm-name "my-alarm" \  
  --metric-name "CPUUtilization" \  
  --namespace "AWS/RDS" \  
  --statistic "Average" \  
  --period 300 \  
  --evaluation-periods 3 \  
  --threshold 90.0 \  
  --comparison-operator "GreaterThanOrEqualToThreshold" \  
  --dimensions "Name=DBInstanceIdentifier,Value=my-db-instance" \  
  --alarm-actions "<arn of sns resource>"
```

This will monitor the DB instance during a period of 300 seconds (5 minutes) during 3 evaluation periods: $5 \times 3 = 15$ minutes.

If in the three periods, the average is equal or more than 90%, then the alarm will trigger the SNS resource.

You should subscribe to the SNS resource you create by email or SMS.

Connect Deeper

This work was first published in [Practical AWS](#) course.

Join our online community [FAUN](#) and subscribe to our podcast [The DevOps Fauncast](#).

Visit our publication [The Chief I/O](#), subscribe to the newsletter and get cloud native insights from our contributors.